

# kNN and Regression

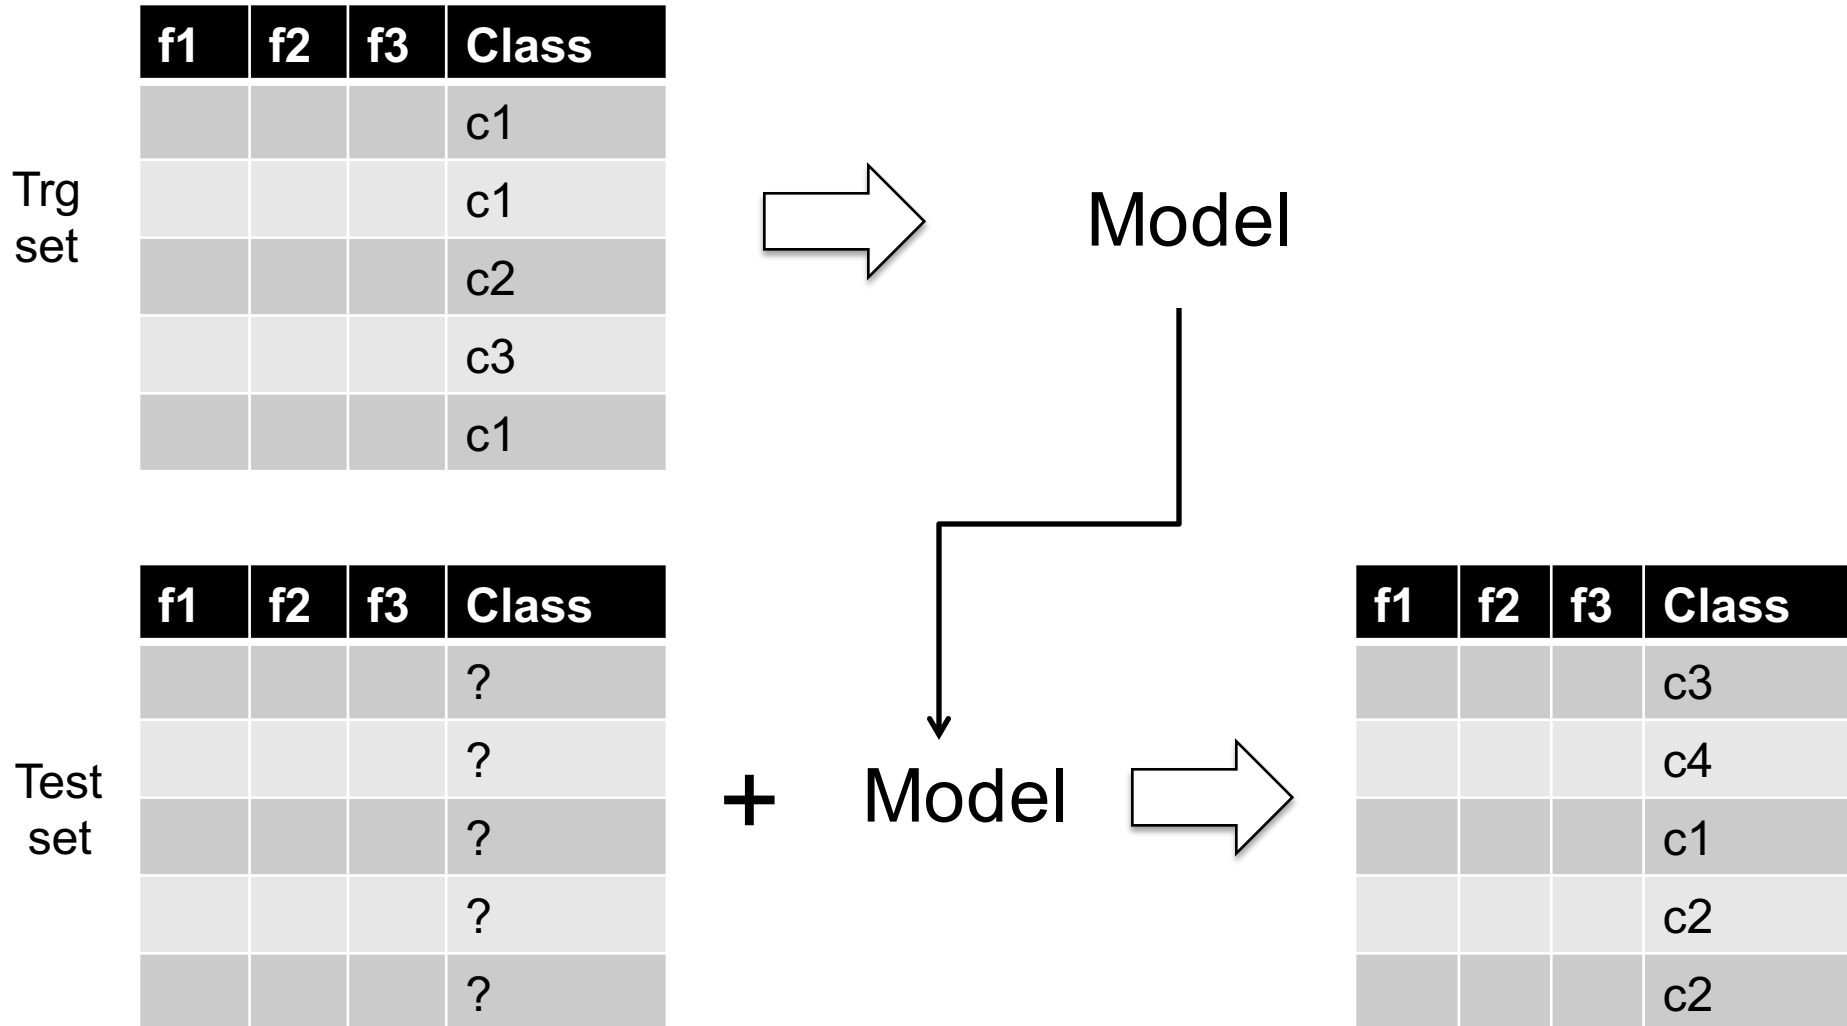
Naren Ramakrishnan



# Recap

- We now know two ML algorithms!
  - Decision Trees and Naïve Bayes
    - Pretty powerful approaches in their own right
    - Form the basis of more powerful methods
      - e.g., Decision Trees -> Random Forests
      - e.g., Naïve Bayes -> Bayesian Networks
- Next
  - We will focus on methods especially suited when features are real-valued
    - Target class is nominal: classification
    - Target class is real-valued: regression

# What's common between DT and NB



# Eager vs Lazy algorithms


- A model-based algorithm is
  - “eager” by definition
  - Most of the hard work is done during the learning phase
- A lazy learner doesn't do much learning
  - Most of the hard work is done during the evaluation/application phase

# Classical lazy algorithm


- Nearest neighbor
  - Also called “instance based learning”
  - Also referred to as a non-parametric method

# Example input


- Consider a real-valued set of attributes
  - Two dimensional



# Basic Idea (k=1)




# Basic Idea ( $k=4$ )





# kNN as a smoother



# How do we choose $k$ ?

- Divide training examples into two sets
  - A training set (80%) and a validation set (say 10%) – test set is separate (another 10%)
- Predict the class labels for validation set by using the examples in training set
- Choose the number of neighbors  $k$  that maximizes the classification accuracy
  - => Cross-validation

# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



$err(1) = 1$

# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



$k = 2$

$err(1) = 3$

$err(2) = 2$



# Leave-One-Out Method

- For  $k = 1, 2, \dots, K$


- $err(k) = 0$

- For  $i = 1, 2, \dots, n$

- \* Predict the class label  $\hat{y}_i$  for  $\mathbf{x}_i$  using the remaining data points

- \*  $err(k) = err(k) + 1$  if  $\hat{y}_i \neq y_i$

- Output  $k^* = \arg \min_{1 \leq k \leq K} err(k)$



$err(1) = 3$


$err(2) = 2$

$err(3) = 6$

# When should we use kNN?



- Advantages
  - Can learn very complex target functions with arbitrary boundaries
  - Fast training time (duh!)
- Disadvantages
  - Can get easily bogged down by noise (e.g., irrelevant attributes)
  - Slow at classification time (duh!)

# Speeding up kNN (kd-tree)





- Each node contains
  - Children information
  - The tightest box that bounds all the data points within the node.

# NN Search by kd-tree





We traverse the tree looking for the nearest neighbor of the query point.

# NN Search by kd-tree





Examine nearby points first: Explore the branch of the tree that is closest to the query point first.

# NN Search by kd-tree





When we reach a leaf node: compute the distance to each point in the node.

# NN Search by kd-tree



When we reach a leaf node: compute the distance to each point in the node.



# NN Search by kd-tree



Then we can backtrack and try the other branch at each node visited.





# NN Search by kd-tree



Using the distance bounds and the bounds of the data below each node, we can prune parts of the tree that could NOT include the nearest neighbor.

# NN Search by kd-tree



Using the distance bounds and the bounds of the data below each node, we can prune parts of the tree that could NOT include the nearest neighbor.


# Curse of dimensionality

- If only 2 attributes (out of say 30) are relevant to the target class, kNN is easily misled by the high-dimensionality of data
- Solution
  - Do feature selection before applying kNN

# Other configurations of kNN

- Instead of choosing top-k nearest neighbors
  - Pick nearest neighbors within some distance
- Other ways of aggregating responses instead of simple majority
  - Weighted k-nearest neighbors

# Linear classifier



# What forms a “model” in a linear classifier?

- *Answer:* The geometry of the line (in 2D) or plane (in higher dimensions)
  - The coefficients of the linear expression of the plane

# Some math

- A hyperplane is typically given by

$$w_1x_1 + w_2x_2, \dots, +w_dx_d + w_0 = 0$$

- Slope of the hyperplane defined by
  - $(w_1, w_2, \dots, w_d)$
- Intercept of the hyperplane defined by
  - $w_0$

# Using a hyperplane as a classifier

- Let

$$g(\mathbf{x}) = w_1x_1 + w_2x_2, \dots, +w_dx_d + w_0$$

- Then

$$y = \text{sign}(g(\mathbf{x})) = \begin{cases} +1 & \text{if } g(\mathbf{x}) \geq 0 \\ -1 & \text{otherwise} \end{cases}$$



# All this sounds good, but..

- How do we learn a classifier?
  - Learning here means “estimating the weights  $\mathbf{w}$ ”
- Solution:
  - Mistake-driven approach
  - Guess all weights
  - Repeat
    - See which examples you make a mistake on
    - Nudge/adjust the weights in the direction to reduce the error

# Does this work in practice?

- Yes, but (a big **BUT**)
  - Only if the examples are linearly separable in the first place
    - Difficult to know this in high dimensions!
- What do you do if examples are not linearly separable?
  - Come up with more complex constructs
    - One hyperplane = one neuron
    - Lots of hyperplanes = Neural network
    - Zillions of hyperplanes = Deep neural network

# What have we learned so far?


- Classifiers
  - Decision Trees
  - Naïve Bayes
  - Nearest neighbors
  - Linear classifiers
- More fancy methods
  - Logistic regression classifier
    - A marriage of Naïve Bayes and linear classifiers
  - Support vector machines (SVMs)
    - Fancier, more robust, linear classifiers

# Regression


- The output/target variable is not a discrete quantity (like PlayTennis) but a continuous outcome
- Many existing methods of classification can be adapted to regression

# Linear regression


- Relationship between TV advertising budget and sales
  - One dependent var and one independent var



# What does it mean to draw a line through the points?



# What constitutes a good fit?




# Ordinary least squares fit

- Find the coefficients of the hyperplane such that
  - Sum of squared errors from predicted to actual values is minimized
  - Also called “linear least squares”
- What is ***ordinary*** about this?
  - The way we measure discrepancy is ordinary
  - Can instead use perpendicular distances!




# Logistic regression

- Suitable for binary regression
  - (e.g., predicting probabilities)



# How logistic regression works



# Logistic regression details

- “Logit” model
  - $\ln[p/(1-p)] = \alpha + \beta X$
- $p$  is the probability that the event occurs
  - $p/(1-p)$  is the “odds ratio”
  - LHS is the log odds ratio
  - RHS is a regular linear expression for a plane

# kNN regression

- Is exactly what it sounds like! 😊
- Combine predictions for nearby points
  - Averaging
  - Interpolation
  - Local linear regression
  - Local weighted regression

# What we have learned thus far

- Lazy methods
  - k-nearest neighbors (both classification and regression)
- Linear methods
  - Again both classification and regression